



Spatial Index Structures

Instructor: Cyrus Shahabi

Outline




- Grid File
- Z-ordering
- Hilbert Curve
- Quad Tree
 - PM
 - PR
- R Tree (next session)
 - R* Tree
 - R+ Tree



Grid File


- Hashing methods for multidimensional points (extension of Extensible hashing)
- Idea: Use a grid to partition the space → each cell is associated with one page
- Two disk access principle (exact match)

Grid File



- Start with one bucket for the whole space.
- Select dividers along each dimension.
Partition space into cells
- Dividers cut all the way.

Grid File




- Each cell corresponds to 1 disk page.
- Many cells can point to the same page.
- Cell directory potentially exponential in the number of dimensions



Grid File Implementation

- Dynamic structure using a grid directory
 - Grid array: a 2 dimensional array with pointers to buckets (this array can be large, disk resident)
 $G(0, \dots, nx-1, 0, \dots, ny-1)$
 - Linear scales: Two 1 dimensional arrays that used to access the grid array (main memory) $X(0, \dots, nx-1)$, $Y(0, \dots, ny-1)$

Example



Grid File Search

- Exact Match Search: at most 2 I/Os assuming linear scales fit in memory.
 - First use liner scales to determine the index into the cell directory
 - access the cell directory to retrieve the bucket address (may cause 1 I/O if cell directory does not fit in memory)
 - access the appropriate bucket (1 I/O)
 - E.g., $X=(0, 1000, 1500, 1750, 1875, 2000)$; $Y=(a, f, k, p, z)$
--- search for $[1980,w]$
- Range Queries:
 - use linear scales to determine the index into the cell directory.
 - Access the cell directory to retrieve the bucket addresses of buckets to visit.
 - Access the buckets.



Grid File Insertions

- Determine the bucket into which insertion must occur.
- If space in bucket, insert.
- Else, split bucket
 - how to choose a good dimension to split?
 - ans: create convex regions for buckets.
- If bucket split causes a cell directory to split do so and adjust linear scales.
- insertion of these new entries potentially requires a complete reorganization of the cell directory---expensive!!!



Grid File Deletions

- Deletions may decrease the space utilization.
Merge buckets
- We need to decide which cells to merge and a merging threshold
- Buddy system and neighbor system
 - A bucket can merge with only one *buddy* in each dimension
 - Merge adjacent regions if the result is a rectangle




Z-ordering

- Basic assumption: Finite precision in the representation of each coordinate, K bits (2^K values)
- The address space is a square (image) and represented as a $2^K \times 2^K$ array
- Each element is called a pixel

Z-ordering

- Impose a linear ordering on the pixels of the image $\rightarrow$ 1 dimensional problem



$$\begin{aligned} Z_A &= \text{shuffle}(x_A, y_A) = \text{shuffle}("01", "11") \\ &= 0111 = (7)_{10} \end{aligned}$$


$$Z_B = \text{shuffle}("01", "01") = 0011$$

Z-ordering

- Given a point (x, y) and the precision K find the pixel for the point and then compute the z-value
- Given a set of points, use a B+-tree to index the z-values
- A range (rectangular) query in 2-d is mapped to a set of ranges in 1-d

Queries

- Find the z-values that contained in the query and then the ranges



$Q_A \rightarrow \text{range } [4, 7]$

$Q_B \rightarrow \text{ranges } [2,3] \text{ and } [8,9]$



Hilbert Curve


- We want points that are close in 2d to be close in the 1d
- Note that in 2d there are 4 neighbors for each point where in 1d only 2.
- Z-curve has some “jumps” that we would like to avoid
- Hilbert curve avoids the jumps : recursive definition

Hilbert Curve- example


- It has been shown that in general Hilbert is better than the other space filling curves for retrieval *
- H_i (order-i) Hilbert curve for $2^i \times 2^i$ array



H_1



H_2





... $H_{(n+1)}$

* H. V. Jagadish: Linear Clustering of Objects with Multiple Attributes. ACM SIGMOD Conference 1990: 332-342

Hilbert Curve- example-2

- Passing through (indexing) all points without crossing itself
- Example: $\langle a, b, c, d, e \rangle \rightarrow \langle 0, 4, 7, 9, 13 \rangle$
H-values
- *Proximity & distance preserving*



$N=2 \rightarrow H: 0-15$

Hilbert Curves

$$H = \ell(P) : [0, 2^N - 1]^d \rightarrow [0, 2^{Nd} - 1]$$



$$d=2: H = \ell(P) : [0, 2^N - 1]^2 \rightarrow [0, 2^{2N} - 1]$$



Quad Trees

- Region Quadtree
 - The blocks are required to be disjoint
 - Have standard sizes (squares whose sides are power of two)
 - At standard locations
 - Based on successive subdivision of image array into four equal-size quadrants
 - If the region does not cover the entire array, subdivide into quadrants, sub-quadrants, etc.
 - A variable resolution data structure

Example of Region Quadtree






PR Quadtree

- PR (Point-Region) quadtree
- Regular decomposition (similar to Region quadtree)
- Independent of the order in which data points are inserted into it
- ☹️: if two points are very close, decomposition can be very deep


Example of PR Quadtree

A node contains $\leq k$ points - here $k=1$

(0,100) (100,100)







Subdivide into quadrants until the two points are located in different regions




PM Quadtree

- PM (Polygonal-Map) quadtree family
 - PM1 quadtree, PM2 quadtree, PM3 quadtree, PMR quadtree, ... etc.
- PM1 quadtree
 - Based on regular decomposition of space
 - Vertex-based implementation
 - Criteria
 - At most one vertex can lie in a region represented by a quadtree leaf
 - If a region contains a vertex, it can contain no partial-edge that does not include that vertex
 - If a region contains no vertices, it can contain at most one partial-edge




PM Quadtree


PM1 quadtree



PM2 quadtree




PM3 quadtree



- Each node in a PM quadtree is a collection of partial edges (and a vertex)
- Each point record has two field (x,y)
- Each partial edge has four field (starting_point, ending_point, left region, right region)

Example of PM1 Quadtree





References

- National Technical University of Athens , Theoretical Computer Science II: Advanced Data Structures
- Jürg Nievergelt, Hans Hinterberger, Kenneth C. Sevcik: The Grid File: An Adaptable, Symmetric Multikey File Structure. ACM Trans. Database Syst. 9(1): 38-71 (1984)
- H. V. Jagadish: Linear Clustering of Objects with Multiple Attributes. ACM SIGMOD Conference 1990: 332-342